

# Sentiment Analysis


## 1. Data Exploration

Our development dataset contains 28754 reviews, while the evaluation dataset has 12323 reviews (42% of the dev).


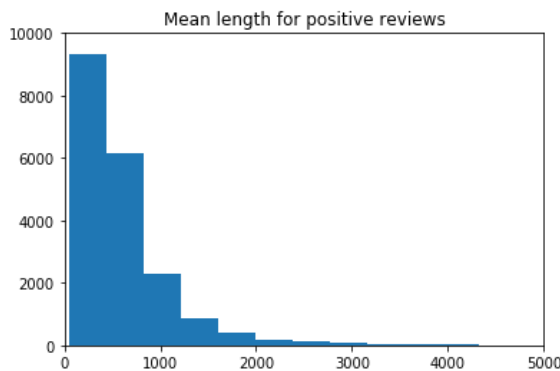
In the development dataset we have 19532 positive reviews (68%) and 9222 negative reviews (32%) → this is an imbalanced dataset.

Concerning outliers, in the development we have 10 reviews of foreign language<sup>1</sup>, while we have 7 in the evaluation.

The following table contains some features with a different distribution among the two classes, so any of these features is a possible candidate.



Features that may have had an impact	Positive	Negative
<b>Text length:</b> mean, stddev	625, 513	864, 740
<b>First name</b> <sup>2</sup> : reviews with at least one first name	26%	54%
<b>#Exclamation marks:</b> mean, stddev	0.94, 2.24	1.63, 3.63
<b>#Question marks:</b> mean, stddev	0.04, 0.32	0.27, 0.98
<b>#Sentences:</b> mean, stddev	6.4, 4.55	7.5, 6.36



## 2. Data preprocessing

### Stopwords

Starting from the `nltk` dataset, I added some words to the list. However, the most important step was to remove negations from the stop-words. The reason is that, given an n-gram, it is possible to have a negation of a certain token. If we leave negations to stop-words, “non è bello” (“it’s not great”) could wrongly become “è bello” (“it’s great”).

### String replacing

I have replaced extremely common words which usually are splitted by a character. For instance, «wi fi» will become «wifi».

<sup>1</sup> Reviews with a number of Italian words < 90% (according to `spacy`)

<sup>2</sup> Italian names from: <https://data.world/axtscz/italian-first-names>

## Tokenization

Possible choices: simple tokenization, lemmatization, stemming. All these features are available in the `spaCy` library. While tokenization simply splits the text into words, lemmatization and stemming usually provide a better pre-processing. Despite some issues (the lemmatization of “la fermata del bus” is “la fermare del bus”), I have decided to use `spaCy` lemmatizer as tokenizer of the `TfidfVectorizer`.

## Outliers elimination

For each review, `spaCy` is able to detect the main language and to assign a score (percentage of words of that language). I have used this feature to detect all the non-Italian reviews and the reviews with a score  $< t_{\text{threshold}}$ . I have translated the reviews without a single sentence in Italian, while I have simply deleted the non-Italian part on the reviews with a score below  $t_{\text{threshold}}$ .

## 3. Algorithm choice

A **Naïve Bayes** classifier, despite the dependence among the tokens [1], offers efficient performances with a reasonable accuracy [2].

A **Random Forest** is a type of classifier which, due to its randomness, is extremely robust to noise and outliers [3]. According to [4], RF could be highly competitive in sentiment analysis with a fine tuning of hyperparameters. Given the structure of a decision tree, it is possible to easily add features without normalize the data.

An **Artificial Neural Network** can simulate the structure of the human brain. As stated by [5], it can be more accurate and precise as compared to NB and SVM algorithms.

A **Support Vector Machine** constructs a set of hyperplanes into a high-dimensional space which split data into classes. It is possible to obtain results comparable to the other algorithms [6]. This classifier allows us to have a specific weight for each class. This could improve our solution given our imbalanced development dataset.

## 4. Tuning and validation

Following scores are performed using `train_test_split` with a test set size = 30% of dataset size.

To tune hyperparameters of vectorizer and classifier I have used `hyperopt`<sup>3</sup>. Once we have defined a distribution, the TPE algorithm [7] will find the optimum by minimizing the objective function, which I have set to  $1 - (p_{\text{pos}} \cdot F_{\text{pos}} + p_{\text{neg}} \cdot F_{\text{neg}})$ , where  $p_{\text{class}}$  is the class percentage and  $F_{\text{class}}$  is the F1 score of *class*.

### 4.1 Data balancing

External dataset used: “Scraping-TripAdvisor-with-Python-2019”<sup>4</sup>.

(pos-neg)	Final distribution (pos-neg)	F1 score variation <sup>5</sup>
Original data	70-30	
Original data – 20%pos → 50-50	50-50	-5%
Original data + (0-100) <sub>external</sub> → 50-50	50-50	-10%
Original data + (50-50) <sub>external</sub>	65-35 <sup>6</sup>	-2.5%
Original data + (70-30) <sub>external</sub>	70-30	<b>+5%</b>

<sup>3</sup> <https://github.com/hyperopt/hyperopt>

<sup>4</sup> <https://github.com/giusepppegambino/Scraping-TripAdvisor-with-Python-2019> > reviewALL.csv

<sup>5</sup> F1 score =  $F1_{\text{pos}} * \%_{\text{pos}} + F1_{\text{neg}} * \%_{\text{neg}}$

<sup>6</sup> Depends on the external dataset size

## 4.2 Vectorizer (TF-IDF)

The values obtained with `hyperopt`, having a space of [0, 0.5] for `min_df` and [0.5, 1] for `max_df`, are those in the following table. The best values have a `max_df` of about 0.3, and a very low `min_df`. The worst results have a `max_df` which varies from 0.5 to 0.95 (so all the available space), but the common factor is the high values of the `min_df`. We can say that we need a low `min_df` (low values could also be zeros) and a `max_df` of about 0.3.

I have used the same approach for the two solutions I have sent. I have obtained opposite results with local trials compared to the leaderboard, so I have uploaded both due to the possibility of having overfitted the test set.

Best values

max_df	min_df	loss
0.346	0.0013	0.0699
0.275	0.0070	0.0733
0.317	0.0009	0.0747
0.292	0.0003	0.0768
0.383	0.0047	0.0776
0.441	0.0002	0.0796
0.314	0.0057	0.0802

max\_df ≈ 0.3  
min\_df: low

Worst values

max_df	min_df	loss
0.503	0.4763	0.4512
0.557	0.2439	0.3378
0.515	0.3939	0.3298
0.524	0.2720	0.3291
0.822	0.4390	0.3098
0.965	0.4975	0.3041
0.958	0.4634	0.2943

max\_df: any  
min\_df > 0.2

## 4.3 Decomposition

Having  $n_1$  of features from the TF-IDF matrix, and reducing the feature size with SVD to  $n_2 \ll n_1$  features, turned out to be worse: we gain 20% of avg F1-score by vectorizing directly using `max_features = n2`.


$n_1$	$n_2$	F1 [neg; pos]
20 000	5 000	0.8425; 0.8710
5 000	None	<b>0.9253; 0.9502</b>

## 4.4 Classifier

	BernoulliNB	GaussianNB	MultinomNB	ComplNB	RandFor(100)	ArtNN	LinSVC
<b>F1 (neg, pos)</b>	0.9279; 0.9518	0.9172; 0.9446	0.9346; 0.9565	0.9385; 0.9581	0.9158; 0.9464	0.9164; 0.9448	<b>0.9379;</b> <b>0.9612</b>
<b>Time (mm:ss)</b>	00:01	00:05	00:00	00:00	01:50	15:00	00:05

ngram_range	F1 [neg; pos]	Time (sec)
(1,1)	0.9106; 0.9498	10
(1,2)	0.9071; 0.9481	15
(1,3)	<b>0.9102; 0.9501</b>	30
(1,4)	0.9066; 0.9478	45
(1,5)	0.9084; 0.9491	60

According to the following ROC curve, if we choose a probabilistic algorithm, the best threshold (the value which maximized the Youden's J statistic<sup>7</sup>) for a review to be positive is 0.681.



## 5. Issues and possible improvements

By comparing the predicted labels of the development dataset with the real ones, it is possible to analyze the wrong predictions. I have noticed that there are two main groups of wrong predictions:

- Approximately neutral: reviews which contains both positive and negative sentences. Possible solution: assign a weight to each sentence (according to adjectives, punctuation, uppercases...). The final prediction would be positive if the summation of positive weights over negatives is above a threshold;
- Too short: reviews containing one or two sentences. Possible solution: use a different classifier with different hyperparameters.

To enhance the algorithm, it could also be possible to remove 1-grams from the tokens if that token will be included in another n-gram which includes a negation.

## References

- [1] P. Gamallo and M. Garcia, "A Naive-Bayes Strategy for Sentiment Analysis on English Tweets," Association for Computational Linguistics, 2014.
- [2] C. Manning, P. Raghadvan and H. Schütze, "Introduction to Information Retrieval," Cambridge University Press, Cambridge, 2008.
- [3] H. H. Parmar and G. Shah, "Experimental and Comparative Analysis of Machine Learning Classifiers," International Journal of Software Engineering and Knowledge Engineering, 2013.
- [4] H. H. Parmar, S. Bhandari and G. Shah, "Sentiment Mining of Movie Reviews using Random Forest with Tuned Hyperparameters," 2014.
- [5] S. J. Livingston, B. S. Selvi, M. Thabeetha, C. Grena and C. S. Jenifer, "A Neural NetworkBased Approach for Sentimental Analysis on Amazon Product Reviews," International Journal of Innovative Technology and Exploring Engineering, 2019.
- [6] A. Muscolino and S. Pagano, "Sentiment analysis, a Support Vector Machine Model based on Social Network data," International Journal of Research in Engineering and Technology, Catania, 2018.
- [7] J. Bergstra, R. Bardenet, Y. Bengio and B. Kégl, "Algorithms for Hyper-Parameter Optimization," 2011.

---

<sup>7</sup> [https://en.wikipedia.org/wiki/Youden%27s\\_J\\_statistic](https://en.wikipedia.org/wiki/Youden%27s_J_statistic)